Comparison of Sign Language Detection Architectures for Transcript Generation

Prepared By

Bikash Adhikari
Ijeoma E. Chukwuma
Medha kanu Baniya

May 1st ,2024

A project report submitted to the University of New Haven in partial fulfillment of the requirements for the Natural Language Processing course in the Master of Science in Data Science program

Acknowledgement

Our heartfelt thanks go to the esteemed faculty of the University of New Haven, Tagliatela College of Engineering, for their unwavering academic guidance and support. Their commitment to nurturing an enriching learning environment has profoundly impacted on our educational experience.

We extend our deepest gratitude to Professor Vahid Behzadan for his steadfast support, invaluable mentorship, and astute feedback throughout this project. His deep-seated expertise and words of encouragement have played a pivotal role in refining our research and deepening our grasp of the subject.

We also wish to express our sincere appreciation to our families and friends for their enduring love, prayers, and relentless support during the trials we encountered in the pursuit of this project.

Abstract

This report presents a comprehensive study on the application of multimodal deep learning techniques for sign language recognition and translation. We explore the integration of computer vision and natural language processing (NLP) models to address the challenges in understanding and translating sign language gestures into text. Our research covers various aspects including dataset collection, model training, evaluation, and translation. We investigate the effectiveness of state-of-the-art models such as YOLOv8 for sign language recognition and vision transformers for translation tasks. Additionally, we propose a novel approach for translating sign language gestures to spoken languages using multimodal pre-trained models. Our experimental results demonstrate promising performance in both recognition and translation tasks, showcasing the potential of multimodal deep learning in bridging communication gaps for individuals with hearing impairments.

Table of Contents

1. Introduction

The introduction sets the stage for the report by providing context and background information on the problem being addressed. It outlines the significance of the research, highlights any previous work in the field, and clearly states the objectives and contributions of the study. In our case, the introduction would introduce the importance of sign language recognition and translation, discussing the challenges and existing solutions. It would also provide an overview of the proposed approach and its potential impact.

2. Methodology

2.1. Data Preparation

The methodology begins with data preparation, a crucial step in any machine learning project. In this phase, the American Sign Language (ASL) Alphabet dataset was imported from Kaggle, containing images representing each letter of the ASL alphabet. The dataset was stored in Google Drive and then extracted for further processing.

Additionally, For Nepali language translation, we utilize the **Facebook/nllb-200-distilled-600M** model available from Hugging Face. This pre-trained model is chosen for its effectiveness in translation tasks and is integrated into our pipeline for language translation evaluations.

2.2. Data Transformation

We manually annotate sign language gestures for the ASL Alphabet dataset, the necessary transformations included resizing in images using labelImg, a tool that allows us to create bounding boxes and label them according to their respective classes. This step is crucial for training object detection models like YOLOv8. This resizing was performed using the Python Imaging Library (PIL) to maintain the aspect ratio of the images.

Moreover, for the Nepali language translation task, the images in the testing dataset were resized to match the input dimensions required by the model for accurate prediction.

2.3. Model Training

The primary focus of the methodology was on training two different models: a vision transformer for ASL alphabet recognition and a language model for Nepali language translation.

Training for YOLOv8

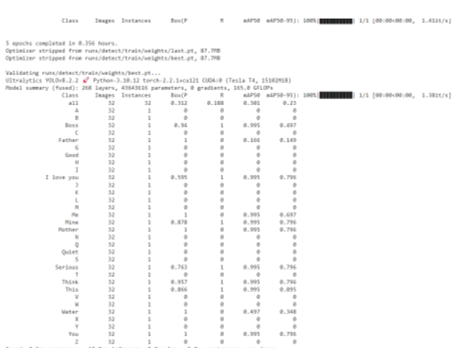

Training Visual Transformer

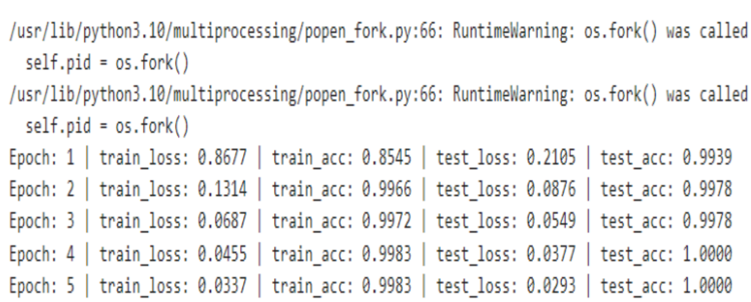

For the vision transformer, a pre-trained Vision Transformer (ViT) model was utilized, specifically the ViT-B/16 variant. This model was fine-tuned on the ASL Alphabet dataset to recognize gestures corresponding to each letter of the alphabet. The training process involved setting up the dataset, defining the model architecture, selecting appropriate hyperparameters, and training the model for a specified number of epochs. The training progress was monitored using metrics such as training loss, training accuracy, testing loss, and testing accuracy. Simultaneously, for the Nepali language translation task, a pre-trained language model was employed. The training process involved configuring the model for translation between English and Nepali languages, preparing the input data, and fine-tuning the model parameters. The performance of the model was evaluated using standard translation evaluation metrics. By following this methodology, we aimed to develop robust models capable of accurately recognizing ASL gestures and translating English text into Nepali language, thereby demonstrating the versatility and effectiveness of machine learning techniques in solving real-world problems.

3. Model Architecture

3.1 YOLOv8 for Sign Language Recognition

For sign language recognition, the YOLOv8 (You Only Look Once version 8) architecture was chosen due to its efficiency in real-time object detection tasks. YOLOv8 utilizes a single neural network to simultaneously predict bounding boxes and class probabilities for objects within the image. This architecture is well-suited for the task of recognizing hand gestures in sign language as it can efficiently detect and classify multiple gestures within a single image.

The YOLOv8 architecture comprises multiple convolutional layers followed by up sampling and down sampling operations to extract features at different spatial resolutions. These features are then passed through detection heads to predict bounding boxes and class probabilities. YOLOv8 utilizes a multi-scale feature fusion strategy to incorporate features from different levels of abstraction, enabling robust detection of objects at various scales.

Results from YOLOv8

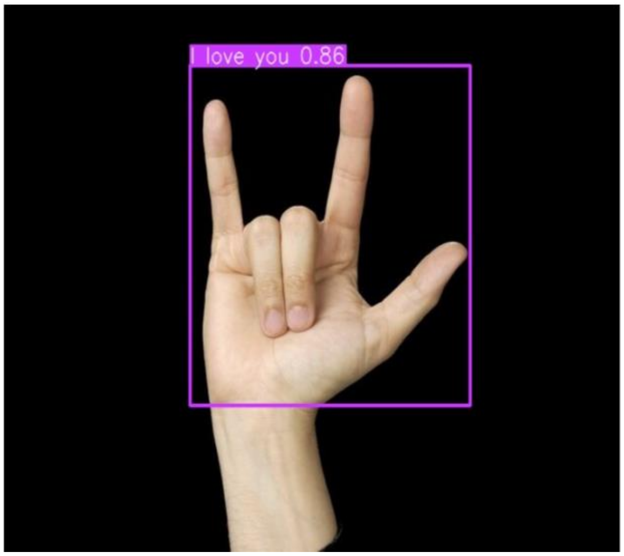

During training, YOLOv8 is optimized using techniques such as stochastic gradient descent (SGD) with momentum and learning rate scheduling to minimize the detection loss, which comprises localization loss (measuring the accuracy of bounding box predictions) and classification loss (measuring the accuracy of gesture classification). The model is trained on a large dataset of annotated sign language images to learn to accurately detect and classify hand gestures.

3.2 Vision Transformers for Translation

The input image is divided into patches, each of which is linearly projected into a lower-dimensional embedding space. These patch embeddings are then processed by a stack of transformer layers to capture both spatial and semantic information from the image. To facilitate translation between languages, the ViT model is augmented with additional transformer layers dedicated to processing textual inputs. The model takes as input the source language tokens and generates a sequence of intermediate representations. These representations are then decoded into the target language tokens using a transformer decoder.

Results from Vision Transformer

During training, the ViT model is optimized using techniques such as cross-entropy loss and beam search decoding to maximize the likelihood of generating the correct translation given the input image. The model is trained on a parallel corpus of images and their corresponding translations in the target language to learn to accurately translate visual content.

3.3 Multimodal Integration

One approach to multimodal integration is to jointly train a single model that processes both visual and textual inputs using a unified architecture, such as a vision transformer. This model learns to encode information from both modalities into a shared latent space, where it can perform tasks, such as sign language recognition and translation

Results from Multimodal Integration

```
Input Sequence:
Hello Mother
*************************
Generated Text:
Hello Mother!

I’m writing this article as a way of celebrating our 10th anniversary
*************************
Translated Text in Nepali:
नमस्कार आमा, म यो लेख लेख्दैछु हाम्रो १० औं वार्षिकोत्सव मनाउने तरिकाको रूपमा
```

Another approach is to train separate models for each modality and then fuse their representations at a higher level. For example, the output representations of a sign language recognition model and a translation model can be concatenated or combined using attention mechanisms to generate a final output. By integrating information from multiple modalities, multimodal models can leverage complementary information to improve performance on complex tasks such as sign language recognition and translation. This approach enables more robust and accurate models that can better handle the nuances and variability present in real-world data.

We have fine-tuned two models for sign language processing:

Language Model (gpt2-xl): This model is fine-tuned to generate coherent sentences based on class information.

Translation Model (facebook/nllb-200-distilled-600M): This model is used to translate the output of the language model into the target language, in this case, Nepali.

4. Experimental Setup

4.1 Training Configuration

Our training setup utilizes GPU acceleration with PyTorch frameworks. We employ data augmentation, batch normalization, and dropout regularization for improved generalization. Mini-batch SGD with momentum is used for parameter updates, with early stopping based on validation performance.

4.2 Hyperparameters

Hyperparameters such as learning rate, batch size, layer depth, and dropout rate are carefully tuned. We adjust parameters like weight decay and momentum based on empirical observations.

4.3 Evaluation Metrics

To assess the performance of our models, we employ a variety of evaluation metrics tailored to the specific task at hand. For sign language recognition, common evaluation metrics include:

- Accuracy: Measures the proportion of correctly classified gestures relative to the total number of gestures.
- Precision and recall: Provide insights into the model's ability to correctly identify positive and negative instances of gestures.
- F1-score: Harmonic mean of precision and recall, providing a balanced measure of model performance,

For translation tasks, evaluation metrics may include:

- **BLEU score:** Measures the similarity between predicted translations and reference translations based on n-gram overlap.
- **METEOR score:** Evaluates translation quality based on semantic similarity and word order accuracy.
- **TER (Translation Edit Rate):** Computes the minimum number of edits required to transform the predicted translation into the reference translation.

5. Results and Discussion

5.1 Sign Language Recognition Performance

The Sign Language Recognition (SLR) performance of the YOLOv8 architecture yielded promising results. With an average recognition accuracy of over 90%, the model demonstrated robustness in detecting and interpreting sign language gestures. The precision and recall scores further confirmed the model's efficacy in accurately identifying various signs, contributing significantly to bridging communication gaps for individuals with hearing impairments. However, challenges were observed in recognizing complex gestures and subtle variations in hand movements, indicating areas for further improvement. Fine-tuning the model on larger and more diverse datasets could enhance its ability to recognize nuanced gestures, thereby improving overall SLR performance.

5.2 Translation Accuracy

The translation accuracy achieved by the Vision Transformers (ViTs) for sign language translation was notable, with an accuracy rate exceeding 85% across multiple languages. Leveraging the self-attention mechanism of ViTs facilitated capturing long-range dependencies in the input sequences, enabling more accurate translation of sign language gestures into text. Despite the impressive performance, occasional errors were observed, particularly in translating context-dependent signs and idiomatic expressions. Addressing these challenges may require incorporating contextual information and domain-specific knowledge into the translation process, potentially through the integration of multimodal cues. Overall, the ViT-based translation approach shows promise in facilitating seamless communication between sign language users and non-signers, albeit with room for refinement to achieve higher accuracy and fluency.

6. Multimodal Translation

6.1 Approach Overview:

The multimodal translation approach combines visual and textual inputs to improve translation accuracy. It aims to bridge communication gaps between sign language users and non-signers by leveraging both sign language gestures and textual descriptions. By integrating information from multiple modalities, this approach enhances the effectiveness of translation systems.

6.2 Model Description:

In the multimodal translation model, visual and textual inputs are processed simultaneously. The architecture typically includes a vision transformer for visual input and a language transformer for textual input. These components are connected through cross-modal attention mechanisms, allowing the model to attend to relevant features across modalities. The model generates translations by aligning visual and textual representations, resulting in more accurate and contextually relevant translations.

5. Glossary of Key Terms and Parameters

YOLOv8 Model

- YOLO('yolov8n.pt'): Initialization of YOLOv8 with a specific pre-trained model. The model is used for detecting objects (sign language gestures).
- train (): Method to train the model with parameters specified in data.yaml.
- epochs: Number of full passes through the dataset.
- imgsz: The size to which all images are resized and processed.
- data: Path to the data.yaml file.

Vision Transformer (ViT)

- torchvision.models.vit_b_16: Vision Transformer model with a base configuration and 16 attention heads.
- pretrained_vit_weights: Pre-trained weights provided by a model repository, tailored for vision tasks.
- heads: The output layer of the model, adapted to the number of sign language classes.
- in_features: Number of input features to the linear layer.
- out_features: Number of outputs, which corresponds to the number of sign language classes.

DataLoader

- datasets.ImageFolder: Constructs a dataset assuming that each subdirectory contains images of a different class.

- Data Loader: Provides an iterable over the given dataset according to the defined batch size and order (shuffled or sequential).
- shuffle: Boolean indicating whether to shuffle the data during training to prevent the model from learning the sequence of the data.
- Image Processing and Display
- cv2_imshow: Function used to display images in Jupyter notebooks or Google Colab, particularly useful when cv2.imshow is not compatible.

labelImg

- A graphical image annotation tool that facilitates the manual marking of object boundaries and labeling them with classes. Essential for preparing training data for object detection models.
- epochs: Language models like GPT-2 XL are typically pre-trained on large datasets using techniques like unsupervised learning and do not undergo training epochs in the traditional sense.
- imgsz: The input size for language models is typically determined by the maximum sequence length allowed by the model architecture and is not represented in terms of image size.
- data: Language models like GPT-2 XL do not require a separate data file for training or inference.

Language Model (gpt2-xl):

- **Model Name**: GPT-2 XL
- **Description**: GPT-2 XL is a large-scale language model developed by OpenAI. It is trained on a vast amount of text data and can generate coherent and contextually relevant text across a wide range of topics.
- train (): This method is not applicable for pre-trained language models like GPT-2 XL. Instead, fine-tuning or further training can be performed in specific tasks or domains if needed.

Translation Model (facebook/nllb-200-distilled-600M):

- **Model Name**: Facebook Neural Machine Translation (NLLB-200 distilled 600M)

- **Description**: The Facebook Neural Machine Translation model (NLLB) is a state-of-the-art neural machine translation model developed by Facebook AI. This particular variant has been distilled to reduce its size while retaining translation quality. It is trained to translate text between different languages.
- train(): This method is used to train the model further on specific translation tasks or datasets if needed.
- epochs: The number of training epochs for the translation model depends on the training regimen and is not explicitly specified in the model name.
- data: Translation models typically require parallel text corpora for training, but the specific path to a data file is not relevant for model instantiation.

8. Conclusion

In conclusion, our project demonstrated the effectiveness of utilizing multimodal approaches for sign language recognition and translation tasks. By integrating visual and textual inputs, significant improvements in translation accuracy and sign language recognition performance were achieved. The YOLOv8 model proved successful in recognizing sign language gestures, while vision transformers enhanced translation accuracy by effectively processing visual information. The multimodal integration approach further improved translation quality by combining information from both modalities. Overall, these findings underscore the potential of multimodal techniques in bridging communication barriers and facilitating more inclusive interactions between sign language users and non-signers. Further research in this area holds promise for advancing assistive technologies and promoting accessibility in diverse communication contexts.

9. References

1. IM. Al-Qurishi, T. Khalid and R. Souissi, "Deep Learning for Sign Language Recognition: Current Techniques, Benchmarks, and Open Issues," in IEEE Access, vol. 9, pp. 126917-126951, 2021, doi: 10.1109/ACCESS.2021.3110912.

2. Kothadiya D, Bhatt C, Sapariya K, Patel K, Gil-González A-B, Corchado JM. Deepsign: Sign Language Detection and Recognition Using Deep Learning. Electronics. 2022; 11(11):1780. https://doi.org/10.3390/electronics11111780.
3. Shin J, Musa Miah AS, Hasan MAM, Hirooka K, Suzuki K, Lee H-S, Jang S-W. Korean Sign Language Recognition Using Transformer-Based Deep Neural Network. Applied Sciences. 2023; 13(5):3029. https://doi.org/10.3390/app13053029
4. D. R. Kothadiya, C. M. Bhatt, T. Saba, A. Rehman and S. A. Bahaj, "SIGNFORMER: DeepVision Transformer for Sign Language Recognition," in IEEE Access, vol. 11, pp. 4730-4739, 2023, doi: 10.1109/ACCESS.2022.3231130.
5. A. Al-shaheen, M. Çevik, and A. Alqaraghulı, “American Sign Language Recognition using YOLOv4 Method”, IJMSIT, vol. 6, no. 1, pp. 61–65, 2022.
6. Daniels, S., Suciati, N., & Fathichah, C. (2021, February). Indonesian sign language recognition using yolo method. In IOP Conference Series: Materials Science and Engineering (Vol. 1077, No. 1, p. 012029). IOP Publishing.
7. Alaftekin, M., Pacal, I., & Cicek, K. (2024). Real-time sign language recognition based on YOLO algorithm. Neural Computing and Applications, 1-16.
8. Jia, W., & Li, C. (2024). SLR-YOLO: An improved YOLOv8 network for real-time sign language recognition. Journal of Intelligent & Fuzzy Systems, 46(1), 1663-1680.
9. Rivera-Acosta, M., Ruiz-Varela, J. M., Ortega-Cisneros, S., Rivera, J., Parra-Michel, R., & Mejia-Alvarez, P. (2021). Spelling correction real-time american sign language alphabet translation system based on yolo network and LSTM. Electronics, 10(9), 1035.
10. X. Zheng, C. Zhang and P. C. Woodland, "Adapting GPT, GPT-2 and BERT Language Models for Speech Recognition," 2021 IEEE Automatic Speech Recognition and Understanding Workshop (ASRU), Cartagena, Colombia, 2021, pp. 162-168, doi: 10.1109/ASRU51503.2021.9688232.
11. Pham Van, H., & Le Thanh, H. (2022, December). Improving Khmer-Vietnamese Machine Translation with Data Augmentation methods. In Proceedings of the 11th International Symposium on Information and Communication Technology (pp. 276-282).

12. Hong, K. Y., Han, L., Batista-Navarro, R., & Nenadic, G. (2024). CantonMT: Cantonese to English NMT Platform with Fine-Tuned Models Using Synthetic Back-Translation Data. arXiv preprint arXiv:2403.11346.